

Learning in AI Planning

Dra. M^a Dolores Rodríguez Moreno

Objectives

Specific Objectives

- Understand what ML is
- Know main ML techniques in AI Planning

Source

- Daniel Borrajo. Slides. Aprendizaje automático en Planificación automática.
- Stuart Russell & Peter Norvig (2009). Artificial Intelligence: A Modern Approach. Chapter 3, 10. (3rd Edition). Ed. Pearsons

Outline

- ML
- Learning agents
- ML in planning
- Learning by macro-operators
- Explanation-Based Learning (EBL)
- Learning by Analogy
- Portfolios
- Conclusions

Definition

ML definition

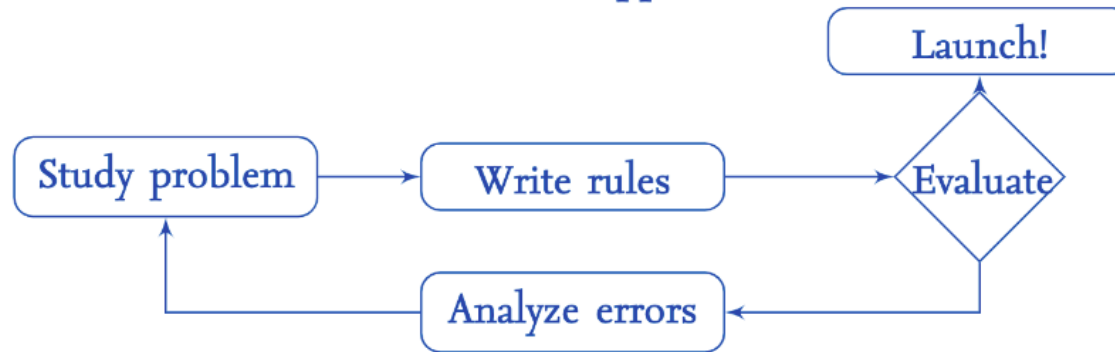
ML is the science (and art) of programming computers so they can learn from data
A. Géron, 2017

Alternative definitions


- Machine Learning is the field of study that gives computers the ability to learn without being explicitly programmed. Arthur Samuel, 1959
- A computer program is said to learn from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience. E. Tom Mitchell, 1997

Definition

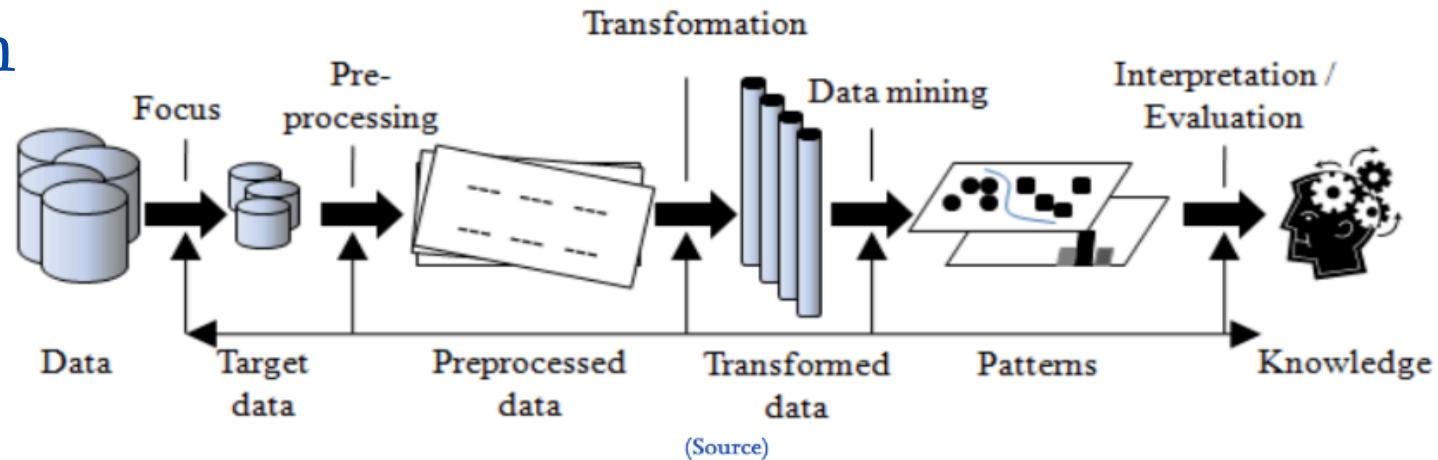
Traditional approach



ML approach



Definition



Steps in any ML application:

1. Data acquisition
2. Selection, cleaning and transformation
3. Machine Learning
4. Learning evaluation
5. Exploitation

The goal in ML is to get a representation of those patterns

Outline

- ML
- **Learning agents**
- ML in planning
- Learning by macro-operators
- Explanation-Based Learning (EBL)
- Learning by Analogy
- Portfolios
- Conclusions

Purpose

- Learning is essential for unknown environments
 - i.e., when designer lacks omniscience
- Learning is useful as a system construction method
 - i.e., expose the agent to reality rather than trying to write it down
- Learning modifies the agent's decision mechanisms to improve performance

Elements

- Design of a learning element is affected by
 - Which components of the performance element are to be learned
 - What feedback is available to learn these components
 - What representation is used for the components
- Type of feedback:
 - **Supervised learning**: correct answers for each example
 - **Unsupervised learning**: correct answers not given
 - **Reinforcement learning**: occasional rewards

Classification

- Inductive L: inferring new knowledge from observations (not guaranteed correct)
 - Concept/Classification L: identify characteristics of class members (e.g., what makes a CS class fun, what makes a customer buy, etc.)
 - Unsupervised L: examine data to infer new characteristics (e.g., break chemicals into similar groups, infer new mathematical rule, etc.)
 - Reinforcement L: learn appropriate moves to achieve goals (e.g., win a game of Checkers, perform a robot task, etc.)
- Deductive L: recombine existing knowledge to effectively solve problems
- Analytical L: reuse the acquired knowledge when presented with similar problems in the future (avoid bad decisions)

Outline

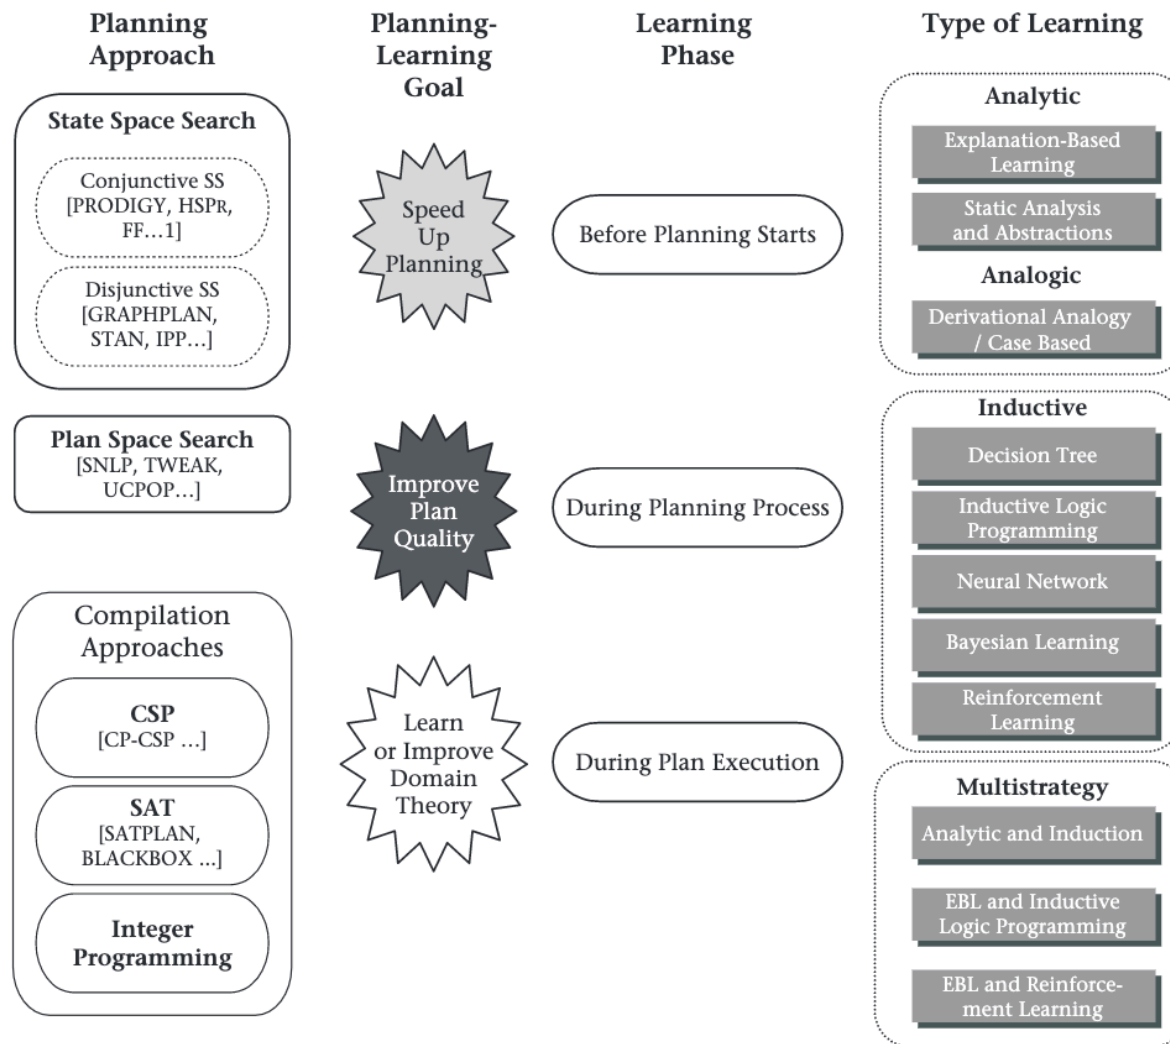
- ML
- Learning agents
- **ML in planning**
- Learning by macro-operators
- Explanation-Based Learning (EBL)
- Learning by Analogy
- Portfolios
- Conclusions

ML in Planning (I)

- Planners efficiency have speed up
- Where can we use learning?

ML in Planning (II)

- Inductive methods
 - Need lots of examples
 - Extract a general description of a *concept*, that is, learn general heuristics
- Deductive/Analytics methods
 - Learn correct knowledge
 - Explain and analyse an example, then learn very specific heuristics
 - Generalize instantiated explanation to apply to other instances
 - Need correct and complete theories
- Analogy methods: CBR



[Source](#)

Outline

- ML
- Learning agents
- ML in planning
- **Learning by macro-operators**
- Explanation-Based Learning (EBL)
- Learning by Analogy
- Portfolios
- Conclusions

Macro-operators (I)

- Building an action model from scratch is a task that is exceedingly difficult and time-consuming even for domain experts
- Some approaches have been explored to learn action models from examples
- A common feature of these works is that they require states just before or after each action to be known
- Statistical and logical inferences can then be made to learn the actions' preconditions and effects

Macro-operators (II)


- First idea to apply learning to planning
- Learning started being applied to the planner STRIPS
- Originally conceived for two-fold purpose:
 - Learning sequences of actions
 - Monitoring execution of plans
- Key idea: create new operators by joining the descriptions of the individual operators that form a plan
- Creation of macro-operators through triangle tables

Macro-operators (III)


- Reuse of past experience (+)
- Re-planning from failures (+)
- Less search depth (+)
- Side-effect: learning operators sub-sequences (+)
- Increased branching factor (-)
- Need to consider utility (-)
 - More we learn less efficient

Macro-operators: example

- Plan: UNSTACK(C,B), PUTDOWN(C), PICKUP(A), STACK(A,B)
- Macro-op
 - Preconditions: on(C,B), on(A,TABLE)
 - Effects: on(A,B), on(C,TABLE)
 - 4 actions
- Generalised by replacing instances for variables



Initial State



Goal

Outline

- ML
- Learning agents
- ML in planning
- Learning by macro-operators
- **Explanation-Based Learning (EBL)**
- Learning by Analogy
- Portfolios
- Conclusions

EBL

- By understanding why an example is a member of a concept, can learn the essential properties of the concept
- Trade-off the need to collect many examples for the ability to “explain” single examples (a “domain” theory)
- Given
 - Goal (e.g., some predicate calculus statement)
 - Situation Description (facts)
 - Domain Theory (inference rules)
 - Operationality Criterion
 - Use problem solver to justify, using the rules, the goal in terms of the facts.
- Generalize the justification as much as possible.
- The operationality criterion states which other terms can appear in the generalized result

EBL

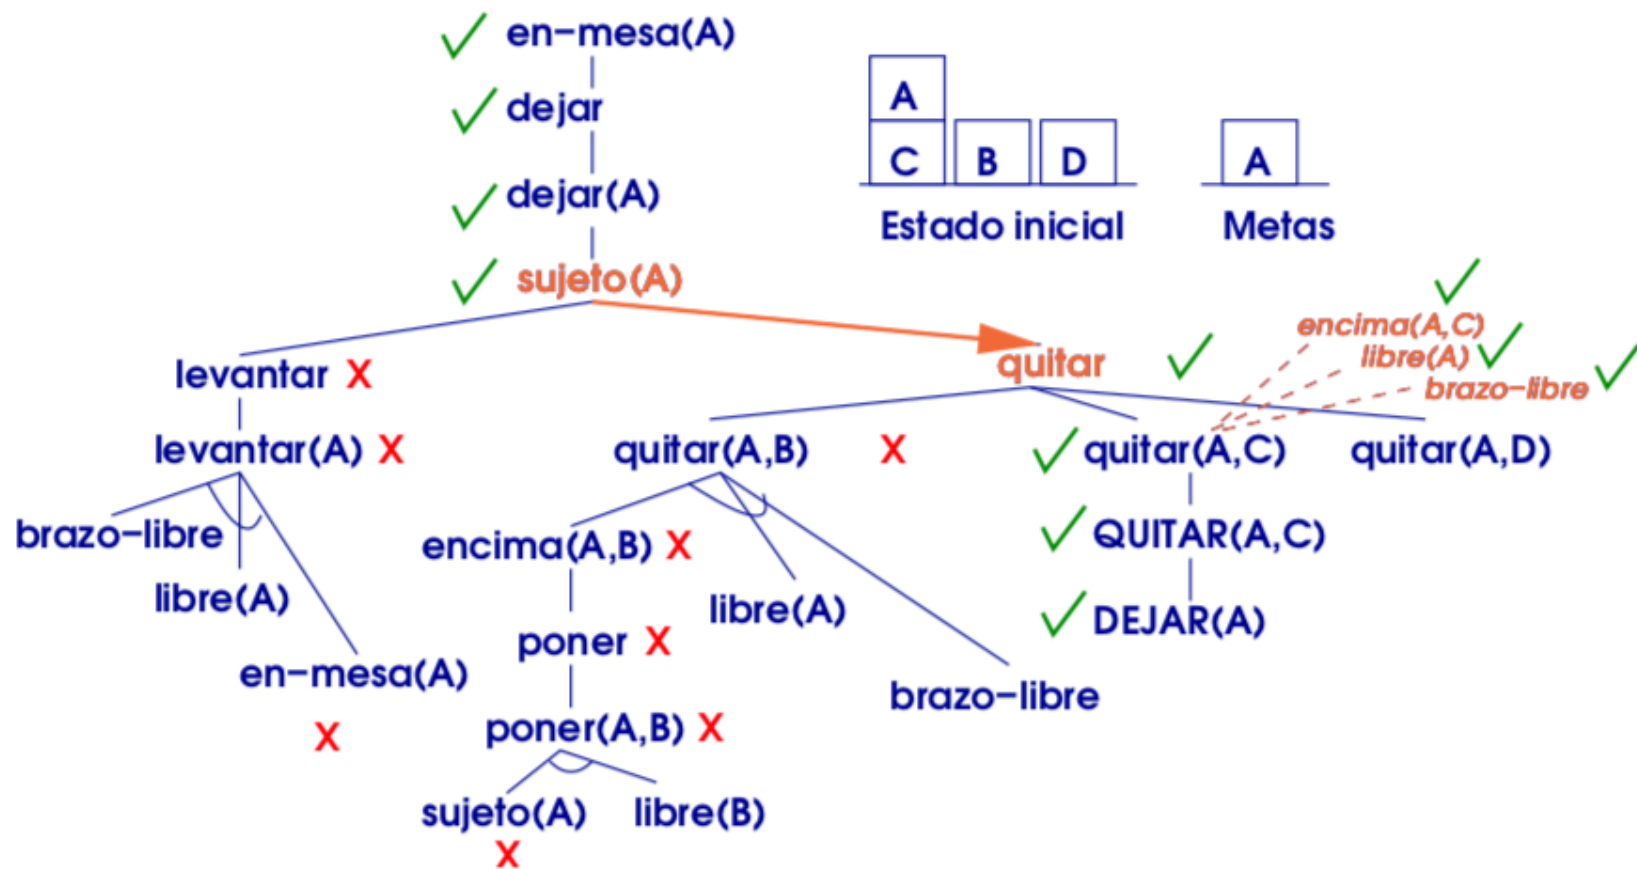
- Actual purpose: re-express target concept in a more operational manner (= efficiency)
- Inputs:
 - Target concept definition: decision to be made
 - Training example: the search episode with its successes and failures
 - Domain theory:
 - Operators used in the search and
 - Objects and possibly relationships in the world which may be used to build the explanation
 - Operationality criterion:
 - Describe concept using terms that are interpretable (efficiently) by the problem solver
 - Several possible criteria

EBL

- Utility problem: the more rules learned, the slower the deliberation

Possible solutions:

- Perform utility analysis and discard low-utility rules
 - Heuristics to learn only effective knowledge
 - Incremental refinement of learned rules
- Factors influencing utility of control knowledge
 - Matching cost (cost of utilization)
 - Frequency of application
 - Savings every time it is applied



Outline

- ML
- Learning agents
- ML in planning
- Learning by macro-operators
- Explanation-Based Learning (EBL)
- **Learning by Analogy**
- Portfolios
- Conclusions

Analogy

- Can apply instance-based learning even when $X \neq R^n$
- \rightarrow need different “distance” metric
- Case-Based Reasoning is instance-based learning applied to instances with symbolic logic descriptions

Analogy

- Humans solve problems combining background knowledge and past experience
- Databases: but they lack generality
- Case-based reasoning: past and new problems need only to be similar for reuse
- PRODIGY/ANALOGY [Veloso, 1994b]

Outline

- ML
- Learning agents
- ML in planning
- Learning by macro-operators
- Explanation-Based Learning (EBL)
- Learning by Analogy
- **Portfolios**
- Conclusions

Portfolios

- Combine a set of planners with time to find the solution
- Autotuning: Determine the parameters that works in some domains
- Types:
 - Problem dependent: IBaCoP
 - Domain dependent: PbP
 - Domain independent: FDSS

Outline

- ML
- Learning agents
- ML in planning
- Learning by macro-operators
- Explanation-Based Learning (EBL)
- Learning by Analogy
- Portfolios
- **Conclusions**

Conclusions

- Learning needed for unknown environments, lazy designers
- Learning agent = performance element + learning element
- For supervised learning, the aim is to find a simple hypothesis approximately consistent with training examples
- Learning performance = prediction accuracy measured on test set